

MOHIT CHAUDHARY

Senior AI/ML & MLOps Engineer | Generative AI • Kubernetes • AWS • Azure Faridabad, India •
+91-9560617079 • mohitsendifier@gmail.com
LinkedIn: [linkedin.com/in/MohitChaudhary](https://www.linkedin.com/in/MohitChaudhary) • GitHub: github.com/mohitsendifier

PROFESSIONAL SUMMARY

Results-driven Senior AI/ML & MLOps Engineer with 7+ years of experience building cloud-native platforms, automating large-scale infrastructure, and deploying enterprise AI/ML solutions across AWS and Azure environments. Over 5 years of hands-on experience supporting AI/ML workflows including model deployment, MLOps automation, data engineering pipelines, Kubernetes-based inference systems, LLM integrations, and AI platform modernization.

Experienced in designing scalable ML infrastructure, deploying Generative AI and RAG-based applications, implementing CI/CD for AI workloads, and enabling real-time model monitoring and governance. Proven expertise in Kubernetes, Terraform, Python, Azure ML, Databricks, MLflow, Docker, and cloud-native AI systems with strong focus on scalability, reliability, security, and cost optimization.

CORE COMPETENCIES

- AI/ML Platform Engineering
 - MLOps & LLMops
 - Generative AI & RAG Pipelines
 - Kubernetes & Docker
 - AWS & Azure Cloud Platforms
 - CI/CD for ML Systems
 - MLflow & Kubeflow
 - Azure ML & Databricks
 - Terraform & Infrastructure as Code
 - LangChain & Vector Databases
 - FastAPI & Python APIs
 - Monitoring & Observability
 - Real-Time AI Inference Systems
 - DevSecOps & Cloud Security
 - ETL & Data Engineering
-

PROFESSIONAL EXPERIENCE

TCS Limited | Gurgaon, India

Systems Engineer → Sr. Systems Engineer → Lead AI/ML & MLOps Engineer

Jan 2019 – Present

Project 1 — Enterprise AI & Data Platform Modernization (Azure)

Client: Multinational Airline

Role: Senior AI/ML Platform Engineer

Jun 2023 – Present

- Designed and implemented enterprise-scale AI/ML platform architecture on Azure supporting ML workloads, real-time inference APIs, ETL systems, and Generative AI services across 12+ subscriptions and 4 business units.
 - Built automated MLOps pipelines using Azure DevOps, GitHub Actions, and Azure ML for training, validation, deployment, rollback, and monitoring of ML models.
 - Deployed containerized ML inference services on AKS using Helm and Terraform supporting 2M+ daily prediction and analytics events with 99.95% platform availability.
 - Developed scalable RAG-based internal AI assistant using LangChain, Azure OpenAI, vector embeddings, and Azure Cognitive Search to improve enterprise document search and operational knowledge access.
 - Integrated Databricks, Azure Data Factory, Kafka, and ML workflows enabling near real-time AI analytics and batch inference pipelines.
 - Implemented ML model observability using Prometheus, Grafana, Azure Monitor, and Application Insights reducing incident detection time by 45%.
 - Automated GPU-enabled inference infrastructure provisioning using Terraform and Kubernetes autoscaling improving AI workload efficiency by 40%.
 - Implemented secure AI governance using Azure Active Directory, RBAC, Key Vault, secret rotation policies, and compliance automation achieving ISO 27001 aligned deployments.
 - Reduced Azure cloud costs by 35% (~₹48 lakhs annually) using workload optimization, spot instances, autoscaling, and automated environment scheduling.
-

Project 2 — Digital Banking AI & Analytics Platform (AWS)

Client: Institutional Banking Organization

Role: Lead MLOps & Cloud Engineer

Jan 2022 – Jun 2023

- Architected end-to-end MLOps and CI/CD pipelines for 40+ microservices and AI-enabled services using Jenkins, GitHub Actions, MLflow, Docker, and Kubernetes.
 - Designed scalable AWS EKS infrastructure for AI/ML model serving, feature engineering workloads, and real-time inference APIs supporting 99.99% uptime across multi-AZ deployments.
 - Implemented ML model deployment workflows with blue-green and canary release strategies using ArgoCD and Kubernetes reducing production rollback incidents by 60%.
 - Developed reusable Terraform modules for AI-ready cloud infrastructure including GPU-enabled compute nodes, S3 data lakes, Aurora, Redis caching, ALB, and secure networking.
 - Supported fraud analytics and predictive banking workloads by integrating ML pipelines with Kafka streaming and automated ETL orchestration.
 - Built centralized AI monitoring stack using Prometheus, Grafana, ELK Stack, Datadog, and PagerDuty reducing MTTR by 40%.
 - Integrated Snyk, SonarQube, AWS Security Hub, and vulnerability scanning into AI delivery pipelines improving compliance posture and reducing critical vulnerabilities by 65%.
 - Containerized legacy analytics services and deployed scalable inference APIs using FastAPI, Docker, and EKS.
 - Collaborated with data science teams to operationalize ML workloads and automate deployment lifecycle management.
-

Project 3 — AI-Enabled E-Commerce Analytics & Cloud Migration

Client: Healthcare & E-Commerce Domain

Role: DevOps & AI Infrastructure Engineer

Jan 2019 – Jan 2022

- Led migration of 50+ enterprise workloads from on-premises systems to AWS enabling scalable AI-ready infrastructure for analytics and recommendation systems.
- Built automated CI/CD pipelines for Python, Java, and PHP applications integrating testing, security scans, and deployment automation reducing release cycles from 2 weeks to 2 days.
- Developed Docker-based environments for analytics and AI experimentation ensuring environment consistency across development and production systems.
- Implemented centralized data ingestion pipelines and cloud storage architecture supporting recommendation engines and reporting workloads.
- Automated provisioning of 200+ EC2 instances using Terraform and Ansible reducing infrastructure provisioning effort by 80%.

- Configured AWS CloudFront, WAF, Route 53, and CDN optimization improving application performance by 55% during high-traffic campaigns.
 - Established backup, disaster recovery, and cross-region replication strategies achieving RPO < 1 hour and RTO < 2 hours.
 - Supported deployment and scaling of AI-enabled analytics services and customer behavior dashboards.
-

AI/ML PROJECTS

Enterprise Generative AI Knowledge Assistant

- Built an enterprise chatbot using Azure OpenAI, LangChain, vector databases, and Retrieval-Augmented Generation (RAG).
- Implemented document ingestion pipelines supporting PDFs, internal documentation, and operational knowledge repositories.
- Deployed scalable inference APIs on Kubernetes with autoscaling and monitoring.
- Improved internal search efficiency and reduced operational support dependency.

Real-Time Fraud Detection Platform

- Supported deployment infrastructure for fraud detection ML models using Kafka, FastAPI, Docker, and Kubernetes.
- Automated ML deployment lifecycle using GitHub Actions and MLflow.
- Enabled real-time inference pipelines for transaction risk scoring.

AI-Powered Recommendation Engine Infrastructure

- Designed scalable AWS infrastructure for recommendation workloads and batch inference processing.
 - Implemented feature engineering pipelines and centralized monitoring.
 - Optimized cloud infrastructure reducing compute costs while improving processing throughput.
-

TECHNICAL SKILLS

AI / ML / GenAI

Python | Scikit-learn | TensorFlow | PyTorch | Hugging Face | LangChain | OpenAI APIs | Azure OpenAI | MLflow | Kubeflow | FastAPI | RAG Pipelines | Prompt Engineering | Vector Databases (FAISS, Pinecone, ChromaDB)

Cloud Platforms

AWS (EKS, ECS, EC2, Lambda, S3, Aurora, RDS, CloudFront, IAM, CloudWatch, Secrets Manager)

Azure (AKS, Azure ML, Azure OpenAI, Azure DevOps, Data Factory, Databricks, Key Vault, Monitor)

MLOps & DevOps

Jenkins | GitHub Actions | Azure DevOps | GitLab CI | ArgoCD | Docker | Kubernetes | Helm | Terraform | Ansible | CI/CD Automation

Data Engineering & Streaming

Databricks | Kafka | Airflow | ETL Pipelines | Spark | Azure Data Factory

Monitoring & Security

Prometheus | Grafana | ELK Stack | Datadog | PagerDuty | Snyk | SonarQube | AWS Security Hub | Trivy | HashiCorp Vault

Programming & Scripting

Python | Bash | Shell | YAML | JSON | Groovy | PowerShell

EDUCATION

B.Tech. in Computer Science & Engineering

Maharshi Dayanand University (MDU), Rohtak

2014 – 2018 | CGPA: 7.5/10

CERTIFICATIONS

- AWS Certified DevOps Engineer – Professional
 - AWS Certified Solutions Architect – Associate
 - Microsoft Azure DevOps Engineer Expert (AZ-400)
 - Microsoft Azure Developer Associate (AZ-204)
 - HashiCorp Certified: Terraform Associate
 - Certified Kubernetes Administrator (CKA)
 - Azure AI Fundamentals (AI-900)
 - Generative AI with Large Language Models
-

KEY ACHIEVEMENTS

- Reduced CI/CD and ML deployment lifecycle duration by 70% through automation and Kubernetes-based build infrastructure.
- Generated ₹1.5 Cr+ in annual cloud savings using infrastructure optimization, autoscaling, spot instances, and workload scheduling.
- Deployed scalable AI/ML inference infrastructure supporting millions of daily events with high availability and automated observability.
- Mentored junior engineers on Kubernetes, Terraform, MLOps, and cloud-native AI deployment practices.
- Built enterprise-grade RAG and AI automation workflows improving operational productivity and knowledge retrieval efficiency.
- Achieved zero critical production incidents over 18 months through automated rollback strategies, observability, and deployment governance.